

Photon-Ring Retarded-Time Tomography I: The Shiva Effect — Null Spaces, Multi-Geometry Observability, and Bounded Inference of Historical Emissivity Level and Morphology

Hina Dixit and Abhinav Chauhan

Preprint draft — theory and controlled synthetic computation only. No telescope detection, no laboratory result, and no recovery from a resolved real photon ring is claimed anywhere in this manuscript.

Two held-out historical inverse results are reported: `STABLE_BASELINE_INCLUSIVE_EMISSIVITY_LEVEL_RECONSTRUCTION` and `AGGREGATE_RESOLUTION_AWARE_MORPHOLOGY_ERROR_REDUCTION`. Neither is recovery of a historical movie.

- campaign tag `paper-I-campaign-final`; the canonical freeze is pinned at commit `8345068676b15ce8f96a76da9d92b159db215f1d`, which is the accepted artifact commit it was first cut against and is not the commit that built this document
 - manuscript built at commit `82e106eee85c871ab6ec9685d823bce8a3c6de59`
 - registry sha256 `2ba66f0209fe1cdec97b8cf5862494c22fb94704318f9601a7a1f9eb4b783796`
 - canonical artifact freeze: 591 artifacts, `artifacts/CANONICAL_ARTIFACT_FREEZE_V2.json`
 - every number below is registered in `artifacts/manuscript/CLAIM_LEDGER.json` with the artifact, row filter and column it was read from
-

Abstract

Light passing near the photon shell of a Kerr black hole reaches a distant observer along a family of increasingly delayed null geodesics, so one image superposes many source epochs. We ask what of that history is recoverable, and separate three quantities that are routinely conflated: algebraic identifiability of a declared source model, historical reach at a given signal-to-noise ratio, and throughput suppression between orders. The forward operator is matrix-free and whitened, built from per-ray Kerr transfer maps on 12 registered spin–inclination geometries with orders $n = 0, 1, 2$, under a pixel-integrated measurement model whose whitened row carries $\sqrt{d\Omega} \, g^3$; the earlier flat-per-row convention made Fisher information scale with pixel count, and is retired with everything it moved recorded rather than absorbed.

The three quantities separate, and they separate in one particular way, which we name the **Shiva effect**: the enrichment of the source model that creates recoverable history destroys identifiability, and the compactness that makes an epoch well posed makes the direct image exactly blind to it. Historical extension is real, modest, and set by inclination rather than spin — the resolved stack sees deeper than the direct image at 12 of 12 geometries, with median recoverable depth 60, 84 and 144 M at $i = 20, 50$ and 75 degrees and a single value across four spins. Retarded-time diversity supplies more of that reach than spatial remapping: flattening the delays costs 0.98 of the higher orders' entire contribution against 0.57 for transplanting the spatial map. Enriching the declared class from 224 to 1056 dimensions drives the resolved operational-rank fraction down 16.8 percentage points and the smallest determined singular value down five orders of magnitude while the deepest recoverable epoch moves by at most 1 grid step of 4 M, and compact temporal support turns the direct image's old-epoch blindness into 84 identically zero columns rather than a large condition number.

Two sealed held-out historical inverse results follow, each judged against a materiality floor fixed before its bank was drawn.

STABLE_BASELINE_INCLUSIVE_EMISSIVITY_LEVEL_RECONSTRUCTION : on 640 truths committed by hash before an operator existed for them, stacking orders extends the anchored stable span from 48 to 80 M at $\text{SNR}_0 = 100$, a gain of 32 M against a threshold of 8 M, on both prior-free estimators and 4 of 4 source families. It is a statement about the age-local emissivity *level*: 98.4% of the field norm is its spatially constant part. **AGGREGATE_RESOLUTION_AWARE_MORPHOLOGY_ERROR_REDUCTION** : on 60 further held-out truths, scored by whichever measure the state's own resolved label selects and with no state excluded, the resolved stack reduces morphology error against the analytic source by 0.164 and 0.133, lower bounds 0.116 and 0.101. Neither an unresolved second image nor total flux reaches materiality, so the gain is attributable to resolving the orders rather than to the photons they carry. Four qualifications travel with that result rather than after it: **MULTI_FEATURE_RECOVERY_NEGATIVE**, **STABLE_MORPHOLOGY_INTERVAL_NEGATIVE**, **FAMILY_HETEROGENEITY** and **DIRECT_BASELINE_SATURATION_QUALIFICATION**.

Between the two, age-local *structure* under our preregistered standard is a bar not cleared at the reference SNR rather than an effect shown to be zero, and it is cleared at ten times that SNR as a secondary result. Neither result is recovery of a historical movie, and we say so at every place both appear. Throughout we carry a negative control: permuting each order's delays, positions and weights independently, which preserves all three marginals and destroys only their pairing, yields a better-conditioned operator than the physical one at 12 of 12 geometries across 16 frozen seeds. Conditioning is not evidence of physical content, and no conclusion here rests on it.

1. Introduction

A photon emitted from the equatorial plane near a Kerr black hole can reach a distant observer directly, or after one or more near-critical windings around the photon shell. Each additional half-orbit adds delay, so the n -th image of a time-varying source shows

that source as it was further in the past. This is the observation behind every proposal to read accretion history out of a resolved photon ring.

Whether that history can be *recovered* is a different question from whether it is *present*, and the difference is where most of the difficulty lives. Three quantities are easy to conflate:

1. **Algebraic identifiability** — does the forward operator have a trivial null space on the declared source model? This is a property of the pair (operator, model class), and a small enough class is identifiable under almost any operator.
2. **Historical reach** — how far into the past can a localized change in the source be detected at a given signal-to-noise ratio? This is a property of where the operator's *sensitivity* lives in retarded time.
3. **Throughput suppression** — how much fainter is each successive image? This is a property of the transfer coefficients alone and says nothing directly about either of the first two.

This paper measures all three on validated Kerr and Schwarzschild ray maps and shows that they separate. It also documents a measurement-convention defect we found and corrected mid-campaign, because several of the conclusions we would otherwise have reported were artefacts of it.

Running through those measurements is one structural fact, and we name it because it recurs in three independent places and is the reason a photon-ring history cannot be read off a reach statistic. **The same enrichment of the source model that creates recoverable history destroys identifiability, and the same compactness that makes an epoch well posed makes the direct image exactly blind to it.** We call this the *Shiva effect*, after the destroyer who is also the creator, and we mean it literally rather than decoratively: the destruction is an exact null space, not an ill-conditioning. Enriching the declared temporal class from 224 to 1056 dimensions moves the deepest recoverable epoch by at most one grid step while driving the resolved operational-rank fraction down by 16.8 percentage points (section 11); replacing the global temporal factor with compactly supported functions gives the direct image 84 identically zero columns on a class where the mirrored global class reports full rank (section 7.1); and the higher orders supply 38 of the old-epoch structural directions the direct image lacks (section 7.2). Creation and destruction are the same operation seen from two sides, and a claim about historical recovery that quotes only one of them is not wrong so much as incomplete.

1.2 Contributions

1. **A whitened, matrix-free operator built from validated ray maps**, with the pixel-integrated measurement convention stated explicitly and the earlier convention retired in public (section 2, section 13).
2. **The separation of identifiability, historical reach and throughput**, measured rather than argued, on 12 geometries (section 5, section 11).
3. **The Shiva effect**: enrichment that creates history destroys identifiability, and compactness that makes an epoch well posed makes the direct image exactly blind to it — an exact null space, not an ill-conditioning (section 7, section 11).
4. **STABLE_BASELINE_INCLUSIVE_EMISSIVITY_LEVEL_RECONSTRUCTION**, a sealed held-out historical inverse result on the age-local emissivity level, with its level-versus-structure decomposition reported in the same section (section 6).
5. **AGGREGATE_RESOLUTION_AWARE_MORPHOLOGY_ERROR_REDUCTION**, a second sealed held-out historical inverse result, on a measure that changes with what the evaluation grid can resolve and excludes no state, with **MULTI_FEATURE_RECOVERY_NEGATIVE**, **STABLE_MORPHOLOGY_INTERVAL_NEGATIVE**, **FAMILY_HETEROGENEITY** and **DIRECT_BASELINE_SATURATION_QUALIFICATION** reported beside it (section 9).
6. **A preregistered negative result on age-local structure** at the reference SNR, preserved as a bar not cleared rather than an effect shown to be zero (section 8).
7. **A nonphysical negative control that beats the physical operator on conditioning**, which refutes conditioning as evidence of physical content (section 10.1).

8. **A governance record that survives its own defects:** every freeze, gate, amendment and retired bank is committed before the result it governs, and the failures are preserved rather than edited (section 16).

1.1 What is deliberately not claimed

No statement here concerns a real observation. The source model is a declared finite-dimensional class, and no result is extrapolated from full rank on such a class to injectivity on a continuum of source functions — that inference is exactly what section 11 shows to be invalid. No machine-learned reconstruction is performed. No geometry-mismatch or order-leakage study is included. The recoverable depth reported everywhere is a detection statement about a localized unit-norm historical mode, never the largest delay any ray happens to carry. The morphology result of section 9 is an aggregate error reduction on single-feature states; it is not recovery of a historical movie, and the two are kept apart wherever both appear.

2. Forward model and measurement convention

2.1 The operator

Each retained ray p of image order n lands at an equatorial source point (r, ϕ) and carries a delay Δt and a redshift factor g . The observation at observer time t_o is built row by row rather than from an order-wide delay:

$$\begin{aligned} z_{\{n,p\}}(t_o) &= d\Omega_{\{n,p\}} g_{\{n,p\}}^3 j(r_{\{n,p\}}, \phi_{\{n,p\}}, t_o - \Delta t_{\{n,p\}}) + \eta \\ \text{Var}(\eta) &= \sigma_\Omega^2 d\Omega_{\{n,p\}} \end{aligned}$$

The distinction from a delay ladder matters. The pilot geometry's orders span *overlapping* retarded windows, so an image order does not correspond to one source epoch, and $a_n j(t_o - n \tau)$ is only an asymptotic summary. A distributed delay kernel and a delay ladder need not have the same null space, and on this operator they do not: the ladder's null space is fixed by the three discrete lags, while the kernel's is fixed by the whole delay distribution within each window.

Whitening under the declared noise gives the row that every audit consumes:

$$A_{\text{ilde}} = \sqrt{d\Omega} / \sigma_\Omega * g^3 * B(r, \phi, t)$$

2.2 The square root is load-bearing

An earlier revision of this work used $c = g^3$ with a flat per-row noise. That convention makes the Fisher information scale with the *number of rows*: splitting one pixel into k equal-area children carrying the same transfer value multiplies the Gram matrix by k . Adaptive ray counts differ by an order of magnitude across geometry and order, and the lensing bands differ in solid angle by roughly three orders of magnitude, so the convention silently reweighted the bands against one another — handing the thin, deep, high-order bands a far quieter detector per unit sky than the direct image.

The corrected convention is invariant under the same split-and-merge test to $5.4e-15$, against 7 for the retired one, and the check is a registered gate. Every result in this manuscript is computed under the corrected convention; section 13 lists what the correction moved.

Derived arms are linear maps of the same resolved data with their covariance propagated, never separate models with their own noise level:

$$\begin{array}{ll} \text{unresolved image} & y_U = L y_R, \quad C_U = L C_R L^T \\ \text{total flux} & y_F = S y_R, \quad C_F = S C_R S^T \end{array}$$

with $C_R = \sigma_{\Omega}^2 \text{diag}(d\Omega)$. One noise density is fixed once from the direct arm's clean response to the declared reference source and shared by every arm. No arm may choose its own σ ; an arm that did would be measuring its own row count.

2.3 Source classes and the localized historical probe

The registered primary class **C224** is 4 radial cubic B-splines (uniform knots in $\log r$) $\times$ 7 real azimuthal Fourier modes $\times$ 8 global temporal DCT modes = 224 dimensions. Every temporal factor spans the whole history, so C224 answers "can a low-dimensional global model be fitted" and cannot be indexed by epoch.

Epoch-resolved questions use the **localized class**: the same radial and azimuthal factors crossed with one compact Gaussian bump in retarded age of half width 3 M. Its $m = 0$, partition-of-unity contraction is a spatially flat scalar probe, and both are used — the scalar probe for the registered depth and innovation statistics, the full 28-dimensional class where a spatially flat probe would be blind (section 5.2).

3. Geometries, ray maps, and the common age grid

Twelve registered geometries span spin a^* in (0, 0.5, 0.9, 0.98) and inclination in (20, 50, 75) degrees, with orders $n = 0, 1, 2$ and 1536 rays retained per order, stratified in screen azimuth, source radius and delay with each band's total solid angle preserved. Observations are 8 times over 20 M.

At exactly zero spin the pinned tracer packages are unusable — one package's critical-curve parameterization is singular there and the other's default Keplerian branch sets the orbital angular velocity to zero — so an explicitly validated Schwarzschild backend is used, and the orbital orientation at $a^* = 0$ is the $a \rightarrow 0^+$ disk-rotation branch. A Schwarzschild black hole has no preferred spin direction, but the disk still has an orbital orientation, and matching the positive-spin limit keeps the grid continuous.

3.1 Delay summaries that converge

The sampled maximum ray delay is an extreme-value statistic set by whichever single ray sits closest to a band edge. It does not converge under refinement and is not used as a historical depth anywhere in this work. The converging summaries are weighted quantiles under three weightings: solid angle ($d\Omega$), throughput ($d\Omega g^3$), and Fisher ($d\Omega g^6$, the squared whitened row weight). Only the first two are properties of the spacetime and the transfer alone; the third depends on the declared measurement model and is labelled as such.

The common age ceiling is fixed from those source-independent summaries over the *whole* grid before any operator is evaluated:

$$\begin{aligned} A_{\max} &= T_{\text{obs}} + 1.25 \max_{(g,n)} \{Q_{0.999}^{\Omega}, Q_{0.999}^I\} + 2h \\ &= 249.6 \text{ M} \rightarrow 252 \text{ M after rounding up to the age step} \end{aligned}$$

with the maximum, 178.9 M, occurring at `a098_i075` . Because the ceiling sits above the deepest ray any geometry carries, 0 of 1152 depth entries in the whole campaign are right-censored: the depths reported below are measurements, not grid ceilings. For contrast, the largest sampled maximum delay anywhere on the grid is 191.5 M — the statistic we do not use.

4. Correctness

Every mechanical gate is the worst case over all 12 geometries (E3C) or all class–anchor combinations (E3D).

gate	status	measured	threshold
<code>E3C_G2_physical_dense_matrix_free</code>	PASS	0	1e-10
<code>E3C_G3_physical_adjoint</code>	PASS	5.52e-14	1e-08
<code>E3C_G4_physical_resolved_unresolved_mixing</code>	PASS	0	1e-10
<code>E3C_G4b_linear_collapse_covariance_propagation</code>	PASS	0	1e-12
<code>E3C_G6_physical_Gram_monotonicity</code>	PASS	9.38e-17	1e-10
<code>E3C_G6b_resolved_dominates_direct</code>	PASS	7.25e-17	1e-10
<code>E3C_G9w_weight_semantics</code>	PASS	4.38e-16	1e-10
<code>E3C_freeze_raymap_hashes</code>	PASS	0	0
<code>E3C_frozen_grid_invariance</code>	PASS	dims=[224] n_ages=[64]	dims=[224] n_ages=[64]
<code>E3C_H2_registered_statistic_is_an_identity</code>	PASS	12	12
<code>E3D_adjoint</code>	PASS	4.89e-13	1e-08
<code>E3D_dense_smoke_comparison</code>	PASS	0	1e-10
<code>E3D_Gram_monotonicity</code>	PASS	1.14e-16	1e-10
<code>E3D_class_nesting</code>	PASS	2.73e-14	1e-10
<code>E3D_enrichment_does_not_lose_rank</code>	PASS	0	0
<code>G10q_continuum_noise_quadrature_invariance</code>	PASS	5.4e-15	1e-10

`E3C_freeze_raymap_hashes` re-checks every ray map against the digest pinned before evaluation;
`E3C_frozen_grid_invariance` checks that the class dimension and age grid were identical at every geometry.
`E3D_class_nesting` is discussed in section 11.

5. Results

5.1 Historical extension is real, modest, and set by inclination

At the reference SNR the resolved stack sees deeper than the direct image at 12 of 12 geometries, and carries strictly positive historical innovation beyond the direct channel's own 99.9% throughput-weighted age boundary at 12 of 12.




Figure 1. Recoverable depth at the reference SNR is flat across four spins at every inclination and steps sharply with inclination. The resolved stack sees deeper than the direct image at all twelve geometries; the margin is largest where the direct image is shallowest.

Recoverable depth, resolved stack (M):

$a^* \setminus i$	20°	50°	75°	trend in inclination
0.00	60	84	144	nondecreasing
0.50	60	84	144	nondecreasing
0.90	60	84	144	nondecreasing
0.98	60	84	144	nondecreasing
trend in spin	constant	constant	constant	

Recoverable depth, direct image alone (M):

$a^* \setminus i$	20°	50°	75°	trend in inclination
0.00	20	60	140	nondecreasing
0.50	20	60	140	nondecreasing
0.90	20	60	140	nondecreasing
0.98	20	60	140	nondecreasing
trend in spin	constant	constant	constant	

The surface is flat in spin and steep in inclination. At each of the three inclinations the four spins give a single value of the resolved depth (1, 1 and 1 distinct values respectively); the medians are 60, 84 and 144 M for the resolved stack against 20, 60 and 140 M for

the direct image. Whatever sets historical reach in this operator, it is the geometry of the retarded windows as seen from a given viewing angle, not the spin parameter.

Depth is a threshold statement, so we report alongside it the threshold-independent **historical innovation**

$$J_{\text{old}} = \text{integral over } a > A_0, 0.999 \text{ of } \log(1 + I(a)) \, da$$

which integrates information about the localized epoch over exactly the region where the direct channel is essentially absent.

Historical innovation, resolved stack:

$a \backslash i$	20°	50°	75°	trend in inclination
0.00	67.86	40.93	12.36	nonincreasing
0.50	69.94	41.21	15.88	nonincreasing
0.90	70.81	40.44	16.07	nonincreasing
0.98	71.53	40.51	16.22	nonincreasing
trend in spin	nondecreasing	nonmonotone	nondecreasing	

Historical innovation, direct image alone:

$a \backslash i$	20°	50°	75°	trend in inclination
0.00	15.02	8.34	5.07	nonincreasing
0.50	15.28	8.78	4.92	nonincreasing
0.90	15.37	8.24	5.09	nonincreasing
0.98	15.40	7.99	5.24	nonincreasing
trend in spin	nondecreasing	nonmonotone	nonmonotone	

J_{old} falls with inclination while depth rises with it. The two are not in tension: at high inclination the direct channel already reaches far back, so the *additional* history the higher orders supply is a smaller increment even though the absolute reach is larger.

The incremental Gram operator $\text{delta_G} = G_{\text{resolved}} - G_{\text{direct}}$ separates information reweighted inside the direct channel's support from genuinely new historical information beyond it. Its rank runs from 195 to 224 of 224:

$a \backslash i$	20°	50°	75°	trend in inclination
0.00	224	224	195	nonincreasing
0.50	224	224	196	nonincreasing
0.90	224	224	196	nonincreasing
0.98	224	224	197	nonincreasing
trend in spin	constant	constant	nondecreasing	

5.2 A registered statistic that could not answer its own question

The registered mechanism test compares the full and substituted age-information curves through a relative discrepancy D . Applied to the delay-only substitution it is **identically zero at every geometry**, with a maximum over the whole grid of 0.0e+00.

This is an algebraic identity, not a finding. The registered probe is spatially flat, and the delay-only substitution replaces only `source_r` and `source_phi`; the scalar curve therefore *cannot* depend on the substitution. We verified the curves are bitwise identical at all 12 geometries and locked that fact behind a gate, so the zero can never be re-read as support for the delay mechanism. Reporting it as such would have been circular.

The literal values are preserved on the record. Amendment 002 adds the comparison that can discriminate: the age-resolved information matrix on the registered 28-dimensional localized class, each probe normalised to unit L2 norm over the emission region, compared through the log information volume $\sum_k \log(1 + \text{SNR}^2 \lambda_k(a))$.

5.3 Delay diversity supplies more of the reach than spatial remapping

Two substitution arms isolate the candidate mechanisms while holding the measurement model, the noise, the source class and the ray weights fixed. `DELAY_ONLY` keeps each order's physical delays and takes the direct order's spatial map; `SPATIAL_ONLY` keeps each order's spatial map and flattens every delay onto the direct order's.

geometry	D_delay (registered)	D_delay	D_spatial	D_direct	D_delay / D_direct	D_spatial / D_direct
a000_i020	0.0e+00	0.1691	0.3108	0.3147	0.537	0.988
a000_i050	0.0e+00	0.1512	0.2060	0.2307	0.655	0.893
a000_i075	0.0e+00	0.1651	0.3036	0.2294	0.720	1.324
a050_i020	0.0e+00	0.1469	0.3229	0.3276	0.448	0.986
a050_i050	0.0e+00	0.1685	0.2041	0.2475	0.681	0.825
a050_i075	0.0e+00	0.2747	0.3418	0.2663	1.032	1.283
a090_i020	0.0e+00	0.1749	0.3344	0.3400	0.514	0.983
a090_i050	0.0e+00	0.1613	0.2183	0.2688	0.600	0.812
a090_i075	0.0e+00	0.2147	0.3099	0.3074	0.699	1.008
a098_i020	0.0e+00	0.1653	0.3386	0.3457	0.478	0.979
a098_i050	0.0e+00	0.1187	0.2416	0.2984	0.398	0.810
a098_i075	0.0e+00	0.1288	0.3734	0.3379	0.381	1.105

The direct arm's discrepancy is the natural zero point: it is what remains when the higher orders are removed altogether. Normalised by it, flattening the delays costs a median 0.98 of everything the higher orders contribute — essentially all of it — while transplanting the spatial map costs 0.57. Delay-only is the closer of the two substitutions at 12 of 12 geometries.

This is weaker than a single-geometry canary suggested. Under the scalar statistic the delay-only arm looked like an exact reproduction of the full operator. On the localized class it is a median 0.165 relative departure against the direct arm's 0.303. Retarded-time diversity carries the larger share of the historical reach; it does not carry all of it, and spatial remapping is not negligible.

Conditioning is reported beside the discrepancy deliberately: spatial remapping can improve κ_+ without moving the oldest detectable epoch, so reading the mechanism off a depth endpoint alone would credit delay with work that spatial diversity does.

5.4 What the order labels are worth

UNRESOLVED_IMAGE sums the orders into a single image plane and pays the summed noise through $C_U = L C_R L^T$.

geometry	T(unresolved) / T(resolved)	J_old(unresolved) / J_old(resolved)
a000_i020	0.333	0.243
a000_i050	0.714	0.235
a000_i075	0.972	0.439
a050_i020	0.333	0.241
a050_i050	0.714	0.245
a050_i075	0.972	0.381
a090_i020	0.333	0.244
a090_i050	0.714	0.240
a090_i075	0.972	0.382
a098_i020	0.333	0.243
a098_i050	0.714	0.233
a098_i075	0.972	0.393

Depth degrades gracefully — the ratio runs from 0.33 to 0.97, reaching 0.97 at high inclination where the direct channel already dominates the reach. Historical innovation does not: the ratio runs 0.23 to 0.44. Explicit order labels are close to dispensable for *how far back* one can see and load-bearing for *how much* is learned about that history.

5.5 Throughput suppression and sensitivity suppression are different exponents

Comparing orders requires matched support. Each order is sampled at the same fractional position within its own retarded window, giving

$$\text{Gamma_sensitivity_matched} = -0.5 \log(I_{\{n+1\}}^{\text{matched}} / I_n^{\text{matched}})$$

against the throughput exponent $\text{Gamma_amp} = -\log(A_g \text{ ratio})$. The one-half is what makes the two comparable: I is a Fisher information and scales as the square of an amplitude, so half its log ratio is the log ratio of the amplitude that Gamma_amp measures directly. All 19 window fractions are retained at every geometry.

geometry	pair	defined / 19	median	IQR	Gamma_amp	difference	window overlap	disposition
a000_i020	0->1	19	2.614	0.397	4.714	2.100	0.000	undefined
a000_i020	1->2	19	1.831	0.124	3.602	1.771	0.534	supported
a000_i050	0->1	19	2.441	0.583	4.274	1.833	0.115	supported
a000_i050	1->2	19	2.187	0.263	3.887	1.700	0.493	supported
a000_i075	0->1	19	1.979	0.909	3.248	1.269	0.308	supported
a000_i075	1->2	19	2.059	0.131	4.057	1.998	0.578	supported
a050_i020	0->1	19	2.723	0.126	4.681	1.958	0.000	undefined
a050_i020	1->2	19	1.794	0.018	3.501	1.706	0.545	supported
a050_i050	0->1	19	2.486	0.561	4.272	1.786	0.114	supported
a050_i050	1->2	19	2.120	0.157	3.784	1.664	0.521	supported
a050_i075	0->1	19	1.530	0.815	2.870	1.340	0.309	supported
a050_i075	1->2	19	2.569	0.343	4.388	1.819	0.519	supported
a090_i020	0->1	19	2.609	0.259	4.620	2.010	0.000	undefined
a090_i020	1->2	19	1.704	0.112	3.203	1.499	0.545	supported
a090_i050	0->1	19	2.492	0.568	4.249	1.758	0.117	supported
a090_i050	1->2	19	1.950	0.211	3.548	1.599	0.557	supported
a090_i075	0->1	19	1.560	0.884	2.875	1.315	0.309	supported
a090_i075	1->2	19	2.310	0.403	4.251	1.941	0.571	supported
a098_i020	0->1	10	1.909	3.409	4.592	2.683	0.435	undefined
a098_i020	1->2	19	1.682	0.042	3.075	1.394	0.536	supported
a098_i050	0->1	19	2.679	1.197	4.230	1.551	0.139	supported
a098_i050	1->2	19	1.684	0.627	3.356	1.672	0.724	undefined
a098_i075	0->1	19	1.611	1.013	2.872	1.262	0.304	supported
a098_i075	1->2	19	2.230	0.261	4.120	1.889	0.559	supported

5 of 24 cells are marked UNDEFINED_NO_COMMON_MATCHED_SUPPORT and are named rather than dropped:

geometry	pair	reasons
a000_i020	0->1	windows_disjoint
a050_i020	0->1	windows_disjoint
a090_i020	0->1	windows_disjoint
a098_i020	0->1	fractions_without_information, negative_exponent_fractions
a098_i050	1->2	negative_exponent_fractions

The zero-overlap cells are the substantive caveat. At $i = 20$ degrees the $n = 0$ and $n = 1$ retarded windows are disjoint, so sampling both at the same fractional position compares two epochs that share no support at all; the statistic is still computed and reported, but "matched support" is not what it means there. These cells remain in every denominator: the interpretable set is 19 of 24, not 19 of 19.

On the 19 interpretable cells the sensitivity exponent is below the throughput exponent in all 19, with median difference 1.700 and range 1.262 to 1.998. A single scalar attenuation exponent describes the flux and misdescribes the information.

At the canary geometry, the corrected medians are 2.486 (0 to 1) and 2.120 (1 to 2) against throughput exponents 4.27 and 3.78. The same quantities computed independently in the single-geometry canary tables give 2.486 and 2.120. The corrected gap is close to a convention-level identity: whitened rows carry $\sqrt{d0\text{mega}}$ where flux carries $d0\text{mega}$, so information scales linearly in solid angle where throughput scales quadratically, and the exponent difference is approximately half the log solid-angle demagnification. **No asymptotic law is fitted:** $n = 0, 1, 2$ does not determine one.

5.6 The geometry surface

These are twelve deterministic registered geometries, not a sample from a population. No p-value, confidence interval or significance claim appears anywhere in this paper; the surface is reported cell by cell.

Operational rank, resolved stack, of 224:

$a \backslash i$	20°	50°	75°	trend in inclination
0.00	203	201	185	nonincreasing
0.50	203	201	184	nonincreasing
0.90	203	203	189	nonincreasing
0.98	203	205	193	nonmonotone
trend in spin	constant	nondecreasing	nonmonotone	

Operational rank, direct image, of 224:

$a \backslash i$	20°	50°	75°	trend in inclination
0.00	133	153	159	nondecreasing
0.50	134	153	157	nondecreasing
0.90	135	154	157	nondecreasing
0.98	133	153	156	nondecreasing
trend in spin	nonmonotone	nonmonotone	nonincreasing	

Conditioning κ^+ , **resolved stack:**

$a \backslash i$	20°	50°	75°	trend in inclination
0.00	7.04e+05	7.03e+05	2.88e+07	nonmonotone
0.50	7.12e+05	7.04e+05	1.73e+07	nonmonotone
0.90	8.56e+05	6.98e+05	1.36e+07	nonmonotone
0.98	1.19e+06	2.00e+05	7.53e+06	nonmonotone
trend in spin	nondecreasing	nonmonotone	nonincreasing	

Read off the cells rather than summarised: the resolved operational rank is constant at $i = 20$, nondecreasing at $i = 50$, nonmonotone at $i = 75$; nonincreasing at $a = 0.00$, *nonincreasing at $a = 0.50$* , nonincreasing at $a = 0.90$, *nonmonotone at $a = 0.98$* . There is no clean monotone spin trend, and section 10.2 shows that what trend there is in rank is partly confounded with the source domain and must not be read as a pure spacetime effect. The inclination dependence is the robust one, and it runs opposite to depth: higher inclination buys reach and costs identifiable directions.

5.7 Arms at a glance

Medians over the 12 geometries: operational rank and its range, positive condition number, historical innovation J_{old} and recoverable depth T_{rec} in M , both at the reference SNR.

arm	op. rank	range	κ^+	J_{old}	T_{rec}
DIRECT_PHYSICAL	153	133–159	9.00e+08	8.29	60
RESOLVED_PHYSICAL	202	184–205	7.84e+05	40.72	84
UNRESOLVED_IMAGE	178	173–183	9.52e+06	9.66	60
TOTAL_FLUX	13	12–14	9.86e+08	6.90	56
DELAY_ONLY	212	193–216	9.50e+04	40.72	84
SPATIAL_ONLY	178	140–209	9.76e+06	8.33	60
EQUALIZED_ORDER_SENSITIVITY	216	196–222	5.04e+04	306.70	120
PAIRING_DESTROYED	221	217–224	1.80e+04	28.72	72

Two rows in that table are not measurement architectures and must not be read as ones. `EQUALIZED_ORDER_SENSITIVITY` removes the physical attenuation between orders by fiat; its large `J_old` and depth are an oracle bound on what attenuation costs, not an achievable configuration. `PAIRING_DESTROYED` is the nonphysical negative control of section 10.1. `TOTAL_FLUX` is a genuine but deliberately impoverished readout: it collapses all spatial information and reaches a median operational rank of 13 against the resolved stack's 202, and is included as a floor.

6. Held-out reconstruction of age-local emissivity level

Result label: `STABLE_BASELINE_INCLUSIVE_EMISSIVITY_LEVEL_RECONSTRUCTION`.

Section 2 through section 5 are an observability audit: they ask what a measurement can distinguish, and they invert nothing. This is the first of the paper's two held-out historical inverse results, and it is deliberately narrow. One geometry, $a^* = 0.5$ and $i = 50$ degrees. One source class, C224. A single sealed bank of 640 truths, hashed before an operator existed for them, scored once.

6.1 What was sealed, and when

The main-test truths were generated, projected into the class and hashed at registration, under commitment `93608e7a7578fe89`. That commitment records the generator version, the projection rule, the coefficient hash of every exact-in-class truth, the analytic rendering rule for off-grid truths, the regime label and the positivity rule. At the moment it was written no operator had been applied to any of those records, no sufficient statistic had been formed, and no score existed.

Before the main run touched a truth, each record was regenerated from its committed seed stream and checked hash by hash; the bank was checked for overlap against every validation split; membership of the class was measured on coordinates other than the ones the truths were projected on; and every hyperparameter was read from the validation selection rather than chosen. Any of those failing stops the run before a score exists.

6.2 Four regimes, because family shift and representation mismatch are

different questions

An earlier pilot called a bank "in class" when its parameters sat near the resolvable scale of C224. They did not sit *in* C224: the exact projection of those truths onto the class still left a structure-normalized residual of roughly four tenths, so an experiment on that bank measured basis mismatch and reconstruction quality together. Membership is now a property of the truth. An in-span truth is defined as `Q_C x` — sample an analytic family, project it, keep the coefficient vector, and treat the synthesised movie as the truth — so it is in the class at every coordinate rather than on the grid it was projected on. Positivity survives because a constant is itself in the class.

Crossing that against family shift gives four regimes: prior-fit families in and out of the span, and the held-out flare family in and out of the span.

6.3 The result

At the reference `SNR_0 = 100`, tolerance `epsilon = 0.25` and quantile `q = 0.95`, the anchored stable span in M under the prior-free truncated SVD:

regime	direct	resolved	difference
IN_CLASS_ID	48	80	+32
IN_CLASS_00D	48	80	+32
OFF_GRID_00D	48	80	+32
OFF_GRID_ID	0	0	0

against a threshold of 8 M declared before the bank was scored. Ridge, the prior-free confirmatory estimator, gives 32 M independently. 4 of 4 prior-fit families reach the threshold. A paired bootstrap over 10000 resamples whose unit is the truth — every noise draw of a resampled truth travels with it, because draws sharpen one history and do not add histories — puts the old-band normalized error reduction at 0.332 with a lower bound of 0.330, and the absolute reduction at 0.429 with a lower bound of 0.419.

6.4 This is level fidelity, not morphology

The registered metric normalises by the whole age-window norm, and **98.4%** of that norm is the spatially constant part of the field. An error under it is therefore not a statement about morphology. So we split the field with an explicit projector,

$$x = P_{\text{level}} x + P_{\text{structure}} x, \quad P_{\text{structure}} = I - P_{\text{level}}$$

where P_{level} is the orthogonal projection, under the plain Euclidean inner product on grid points that the age-window norms themselves use, onto fields constant in space at each source time. Its range is spanned by the class's own 8 temporal modes rendered as spatially uniform fields, so it is a subspace of the declared class rather than an approximation to one, and it carries the positive baseline every family sits on.

Each component below is normalised by the norm of *that* component of the truth, not by the whole-field norm. The two columns are therefore not two parts of one unit budget and must not be added: the level component holds 98.4% of the norm, so a level error of 0.266 is a much larger absolute error than a structure error of 0.927.

quantity	direct	resolved
level error	0.266	0.028
structure error, all ages	0.927	0.466
structure error, old band	1.169	1.141

The level is recovered and the resolved stack recovers it far better. Structure improves across all ages but remains above one in the old band for both arms, which is to say that neither arm recovers old-age morphology at all. The headline is reconstruction of the age-local emissivity **level** under a baseline-inclusive field metric. It is not detailed old-age movie morphology and nothing in this paper should be read as claiming otherwise.

6.5 The gain is in directions the direct image already sees

Each arm supports its own data-supported subspace, and they are not the same size: 154 directions for the direct channel against 202 for the resolved stack, the extra ones being precisely the weakly determined directions the direct image cannot see. Comparing each arm's error on its own subspace therefore penalises the arm that sees more. Judged like for like on the direct channel's own subspace,

the error falls from 0.986 to 0.579. The gain is not an artefact of adding hard directions to a norm, and it is not a prior effect: both estimators here are prior-free.

6.6 Structure recovery is a separate, much-higher-SNR result

Reported apart from the above because it concerns a different quantity at a different operating point. The onset of nonzero age-local structure recovery is **unchanged** between the arms, at `SNR_0 = 30000`, roughly three orders of magnitude above the reference. What differs is the span at that point: 40 M for the direct image against 76 M for the resolved stack. Resolved higher orders extend the structural historical span once that regime is reached; on this bank they do not lower the SNR at which it begins.

6.7 What the sealed bank does not establish

`OFF_GRID_ID` is a negative result and is preserved as one. The exact projection of those truths clears the tolerance, so the tolerance is reachable there and the reconstructions simply do not reach it: recovery of histories outside the declared class is not established by this campaign.

`OFF_GRID_00D` passes, and its passing is a mild-mismatch diagnostic rather than evidence of off-grid robustness. Its representation floor is far smaller than `OFF_GRID_ID`'s, so it is a weaker test of leaving the class, and it is reported as one.

Uncertainty is withdrawn. The joint calibration of the two probabilistic estimators failed a gate whose acceptance band was frozen in advance, and the declared fallback applies: they are retained as point estimators only. No credible interval, posterior movie or coverage statement appears anywhere in this paper.

Integrity controls: 4000 null-pair controls drawn from a bank hashed before scoring, of which 78 exceed the equal-prior Gaussian Bayes bound, consistent with multiplicity and with no evidence that the pipeline reads information the likelihood does not contain.

7. Compact temporal support, and what it costs the direct image

Sections 5 and 6 use a class whose temporal factor is eight global cosines. Every coefficient there is supported on the whole history, so a fit constrained where rays land also determines the field where none do, and no table can separate measured depth from cosine extrapolation. This section removes the mechanism rather than arguing about it.

We replace the temporal factor with compactly supported degree-one B-splines on dyadic node sets, giving three classes `L224`, `L448` and `L1056` that mirror the dimensions of `C224`, `C448_T` and `C1056_ST` exactly. Nesting is arithmetic, not numerical: the coarse node set is the even half of the fine one. Each temporal function occupies a quarter of the history at `L224` and an eighth at `L448`, against all of it for every cosine mode.

7.1 Exact old-epoch null directions

A coefficient whose support contains no ray gives an identically zero column, so "the direct image cannot see this epoch" becomes a statement about the null space. Measured on the same rays:

class	direct rank	identically zero direct columns	mirrored global class	direct rank
L224	140	84	C224	224
L448	252	196	C448_T	411
L1056	523	517	C1056_ST	911

C224 is full rank on its own global temporal subspace, and that number is correct. It establishes identifiability of the 224 global coefficients and says nothing either way about epoch-local identifiability, which C224 cannot pose because none of its coefficients is confined to an epoch. L224, same dimension over the same rays, can pose it, and answers that 84 directions are invisible to the direct image outright.

7.2 Higher orders supply what the direct image lacks

Restricting to old-epoch temporal functions and projecting out the level component gives the old structural subspace. Operational ranks there:

class	direct	resolved	unresolved image
L224	0	38	21
L448	0	93	46
L1056	0	185	91

The direct image is at zero everywhere, with a largest singular value of $1.8\text{e-}14$ — numerical zero. Orders 1 and 2 convert a subspace the direct image cannot see at all into tens or hundreds of measurable directions, and roughly half of that survives summing the orders into a single image plane.

No truth was drawn and no estimator was fitted for this section. It is a statement about the operator.

8. Does the added support reduce reconstruction error?

Section 7 shows the resolved stack observes directions the direct image cannot. Whether that converts into a materially better reconstruction is a separate question, and this section reports a preregistered test of it that the resolved arm does not pass at the reference SNR.

Truths here are synthesised *exactly* in the class, so the representation floor is zero and the error measured is reconstruction error alone. Three structure-first banks replace the baseline-dominated construction of section 6: two non-negative structure-balanced banks carry the physical claim, and a constant-flux bank is retained as a signed linear stress control. Projection into the class removes structure, so the bank nominally at structure fraction 0.80 realises 0.66, and the constant-flux bank reaches a negative mass fraction of 0.197 — it is not everywhere a non-negative emissivity field and cannot carry a source claim.

The endpoint is the paired relative reduction in old-band structural error against the direct image, aggregated with equal weight over bank-family cells. Materiality was fixed in advance: median $\geq 10\%$, both bootstrap lower bounds $\geq 5\%$, at least three of four families, every bank in scope positive, null controls passing, and both estimators agreeing on the same class.

SNR ₀	estimator	median	median CI low	cell-balanced mean	mean CI low	families
100	TSVD	0.168	0.099	0.116	0.039	3/4
100	ridge	0.178	0.096	0.098	0.000	3/4
1000	TSVD	0.275	0.198	0.259	0.210	4/4
1000	ridge	0.304	0.225	0.286	0.237	4/4

At the registered SNR₀ = 100 the medians clear the 10% bar and their own lower bounds clear 5%, but the cell-balanced means fall to 0.116 and 0.098 with lower bounds of 0.039 and 0.000, below the 5% floor, and both drop to three of four families. **The preregistered materiality standard is not met.** The central estimates are positive; what fails is the bar, and we report the bar rather than the sign.

At SNR₀ = 1000 the same arm meets every criterion on both estimators and all four families. That is a secondary result. A gain at tenfold higher normalised SNR does not substitute for the registered point.

The signed constant-flux bank alone gives 0.358 and 0.399, the largest of the three, and pooling it in is what made an earlier version of this analysis read as material. Excluding it, as the design requires, removes the pass. We report that explicitly because it is the most instructive number here: a structural effect that lives mainly in a signed diagnostic is a linear inverse-problem result, not evidence about emissivity histories.

8.1 No stable structural interval

The companion endpoint asks for a contiguous age interval over which the structural error stays under 25% for 95% of truths. Under both noise semantics — averaging the noise draws per truth, and the stricter joint criterion over truth and noise together — every span is zero at both SNRs, so the resolved-minus-direct difference is 0 M against a threshold of 8 M. With the representation floor at zero, nothing blocked the criterion except the reconstruction itself. **This is a real negative result**, and it is complementary to section 8 rather than contradictory: an average-error improvement is not a recovered movie.

No sealed main was executed and none is authorised. Every sealed commitment is preserved unscored.

9. Held-out morphology, measured at the resolution the grid actually has

Result label: AGGREGATE_RESOLUTION_AWARE_MORPHOLOGY_ERROR_REDUCTION , with MULTI_FEATURE_RECOVERY_NEGATIVE , STABLE_MORPHOLOGY_INTERVAL_NEGATIVE , FAMILY_HETEROGENEITY and DIRECT_BASELINE_SATURATION_QUALIFICATION . The four qualifications are part of the result, not commentary on it, and section 9.4 states each one with the number that supports it.

Everything above measures emissivity *level*, or structural error under a norm that treats every direction alike. Neither says whether the *shape* of the source at a given epoch is recovered — how many features it has, where they are, how bright they are. That question needs a measure that knows what the evaluation grid can resolve, and a preregistered statement of what counts as success before the truths are drawn. This is the second of the paper's two held-out historical inverse results.

9.1 The source objects, before any measurement

An earlier attempt at this question was closed without a result. Its declared source range admitted two-feature configurations its declared evaluation grid could not separate, and the two specifications had never been checked against each other. That is a contract between the source model and the grid, and it is decidable from the source alone.

So it was decided from the source alone. 169 source objects — 168 after excluding the preserved failure canary, which is a regression fixture and never evidence — were classified at every age by topographic prominence on three nested grids, with **no ray map imported and no observation operator constructed**, asserted by a guard before and after. Of 10248 source states, 3375 — 32.9% — carry more than one resolved feature.

The number that matters for what follows is how often projection into a source class destroys a distinction the source itself has. On states both finest grids agree are multi-resolved, two-hotspot sources merge into one feature at a rate of 0.136 in the claim-bearing class and 0.299 in the representation-limited control. Pooling in the states only the finest grid calls multi raises these to 0.262 and 0.356; we report the stratified rates, because a state whose multiplicity the analysis grid cannot settle should not be counted as a merger the physics caused. The narrowest feature the claim-bearing class can represent at all is 2.00 M.

None of this involves an estimator. It is what the source objects are, and it sets the ceiling everything below is measured against.

9.2 A measure that changes with the state, and excludes nothing

Each (truth, age) state carries a reconciled label — resolved single, resolved multiple, blended, dead, or ambiguous between the two finest grids — and the label selects the error measure: exact optimal assignment between feature sets, a blended-descriptor distance, or an amplitude comparison. Each is normalized to its own worst case, so the three are commensurable, and **no state is dropped**. Ambiguous states are scored by the blended measure, which is the conservative choice: a state whose multiplicity the grid cannot settle must not be scored as though it had a definite feature count.

Two targets are reported for every arm, always together. `CLASS_CONDITIONAL` compares against the best in-class projection of the truth and asks what the estimator achieved given the representation. `PHYSICAL_END_TO_END` compares against the analytic source and carries the physical claim: an improvement on the conditional target alone is a statement about estimation, not about the source.

9.3 The held-out result

60 truths were drawn from a bank whose per-family commitments were committed before the draw, with 4 paired noise draws each. The hyperparameters came unchanged from an earlier selection split and the runner has no sweep. Materiality was declared before the bank existed: median ≥ 0.10 and paired bootstrap lower bound ≥ 0.05 .

In the claim-bearing class at $\text{SNR}_0 = 100$, the resolved arm reduces the all-state morphology error against the direct image by 0.164 under ridge and 0.133 under TSVD on the physical target, with lower bounds 0.116 and 0.101. On the class-conditional target the reductions are 0.158 and 0.160, with lower bounds 0.114 and 0.092. All eight quantities clear the declared floor.




Figure 3. The two held-out historical inverse results, on separate axes because they measure different things in different units. Left: STABLE_BASELINE_INCLUSIVE_EMISSIVITY_LEVEL_RECONSTRUCTION, the anchored stable span against the threshold declared in advance. Right: AGGREGATE_RESOLUTION_AWARE_MORPHOLOGY_ERROR_REDUCTION, median reduction with the paired bootstrap lower bound as the bar, against the 0.10 materiality floor. Both estimators, both targets.

Both controls behave. The unresolved-image arm — a second image with the same extra photons but no order labels — reaches +0.000 and +0.015, and the total-flux arm -0.020 and +0.036; neither is material. The benefit is therefore attributable to resolving the photon-ring orders rather than to the additional photons an unresolved second image also carries. In the representation-limited control class the same comparison gives 0.101 and 0.061 — material under one estimator and not the other, which is what a representation limit looks like when it is showing rather than being asserted.

9.4 What this result is not

FAMILY_HETEROGENEITY. It is an average over a heterogeneous set. 5 of 12 family-estimator cells are material, and the simplest source in the bank — one moving hotspot — is negative under both estimators:




Figure 4. FAMILY_HETEROGENEITY. Per-family median reduction with the paired bootstrap lower bound as the bar; a ringed marker is a cell that clears the declared floor. Five of twelve family-estimator cells are material, the simplest source in the bank is negative under both estimators, and ten truths per family makes every interval here wide. The aggregate is an average over this spread, not a uniform gain.

The twelve exact values are in the evidence ledger rather than repeated here. Ten truths per family makes every interval wide, so no per-family claim is supported in either direction. What is supported is that the aggregate is not a uniform improvement and must not be quoted as one.

MULTI_FEATURE_RECOVERY_NEGATIVE . Two-feature recovery does not reach materiality anywhere. On the 17 truths of 60 that carry a stable multi-resolved state at all, the absolute assignment cost for the resolved arm is 1.208 and 1.302, where 1.0 is one whole feature wrong. The measure resolves that a state has two features and recovers the morphology of one. It does not recover the pair, and the sealed bank reproduces on held-out truths exactly the split an earlier validation found.

STABLE_MORPHOLOGY_INTERVAL_NEGATIVE . There is no stable morphology interval. 0 M for every arm, estimator, class and SNR. At the reference SNR the resolved arm does not even extend how far back morphology stays in tolerance: its mean reach is 3.06 M against the direct image's 3.32 M, lower under both estimators. At tenfold SNR the ordering reverses, 3.45 M against 2.92 M — worth stating, because the reference-SNR comparison is the claim-bearing one and reads as a general statement if the SNR is left off. A per-age error reduction is not a recovered history at either.

DIRECT_BASELINE_SATURATION_QUALIFICATION . The baseline is partly floor-limited. 56.6% of direct-image states under ridge and 55.0% under TSVD sit at the measure's ceiling, so the mean substantially counts how many states failed outright rather than how far off the recovered shape was. The absolute error the resolved arm reaches is 0.598 and 0.650, against the direct image's 0.727 and 0.755, on a scale whose worst case is 1.0.

Read together: an aggregate, single-feature morphology error reduction on held-out truths, attributable to order resolution, in one geometry and two source classes. Not accurate recovery of a historical movie, which this campaign does not establish at any point.

9.5 Two integrity qualifications

The authoritative bank-drawing stage executed against a working tree that was not clean on the registered pathspecs: a deterministic-hash repair had to be present before the stage could be re-run, and it was committed together with the hashes that stage produced. The porcelain diff is recorded and hashed in the run manifest, and the scoring stage ran at that commit with a clean tree and reproduced every bank commitment. The attestation is not clean and is not presented as clean.

The repair itself is worth stating, because it is a reproducibility failure mode that recurred three times in this campaign. Python salts string hashing per process unless the seed is fixed before the interpreter starts, so a builtin `hash()` of a label string can never match across two processes. The integrity field built that way reported every held-out truth as tampered. Truth content and seeds were SHA-256 throughout, so nothing drawn was affected, and the commitment gate refused to let scoring proceed. The check now scans every script for a bare builtin `hash()` call.

10. Controls

10.1 A physically meaningless operator is better conditioned than the real one

`PAIRING_DESTROYED` permutes delay, spatial position and quadrature weight independently within each order. All three marginals survive; only their pairing is destroyed. Over 16 frozen seeds per geometry it reaches a median operational rank of 220 and a worst-case conditioning of $5.81\text{e}+04$, against the physical resolved operator's median $7.84\text{e}+05$. At 12 of 12 geometries the *worst* seed is still better conditioned than the physical operator.

This is a nonphysical control and is never ranked as an alternative measurement architecture. Its purpose is to refute a specific inference: any argument that reads good conditioning, high operational rank or a large stable rank as evidence that an operator is capturing physical structure is refuted by these rows.

10.2 The source domain moves with spin, and we measured how much

The registered radial support is geometry-dependent — the knots follow the rays, and a higher-spin geometry's rays reach closer to the horizon. A spin trend in any rank statistic is therefore partly a statement about the source domain. Rather than assume the confound away, we repeated three anchor geometries on one fixed radial interval in r/M with identical knot locations.

Operational rank moves under the common support in 12 of 24 anchor–arm combinations, by at most 5 of 224. Recoverable depth is unchanged in 23 of 24. Read strictly: the spin trends in operational rank are partly confounded with the source domain; the depth and innovation results are not.

Separately, near-horizon *ray* coverage does not imply that the emissivity model physically emits to the horizon. Ray-map support and assumed source emission are different objects and are kept apart throughout.

11. Full rank on a declared class is a statement about that class

The registered class C224 reaches full column rank under 21 of 21 arm–anchor combinations. That fact is often where such analyses stop. It should not be, because it is a joint property of the operator and a 224-dimensional model, and the way to find out how much of it is the model's poverty is to enrich the model.




Figure 2. The Shiva effect, in one geometry family. Enriching the declared temporal class drives the resolved arm's operational-rank fraction down while the deepest recoverable epoch is unchanged. Identifiability and historical reach are different quantities and they move independently.

We do so along a nested ladder on three anchor geometries. The azimuthal and temporal factors are literal prefixes, so their columns are preserved; the radial cubic B-spline factor is refined, so its columns move while its span is contained in the enriched span, to $2.7\text{e-}14$. Rank statements depend on the span, and that is what the nesting gate checks. Exact adjoint, Gram monotonicity and a dense smoke comparison were run at every class including the 1056-dimensional one; the worst dense-versus-matrix-free discrepancy is $0.0\text{e}+00$.

class	radial	azimuthal	temporal	dimension	resolved operational rank	as a fraction	sigma_min+
C224	4	7	8	224	201	0.897	1.63e-02
C448_T	4	7	16	448	365	0.815	7.16e-05
C528_S	6	11	8	528	441	0.835	2.87e-04
C1056_ST	6	11	16	1056	770	0.729	1.33e-07

Rank deficiency by arm and class. Each cell gives how many of the three anchor geometries are rank-deficient and the range of nullities among those that are; `full` means every anchor reaches full column rank. Aggregating to a single worst nullity would report a deficiency for an arm that is clean at two anchors out of three, which is a different finding:

arm	C224 (224)	C448_T (448)	C528_S (528)	C1056_ST (1056)
DIRECT_PHYSICAL	full	3/3, -4 to -100	1/3, -6	3/3, -109 to -265
RESOLVED_PHYSICAL	full	1/3, -1	full	3/3, -11 to -35
UNRESOLVED_IMAGE	full	1/3, -1	full	3/3, -14 to -35
DELAY_ONLY	full	1/3, -1	full	1/3, -36
SPATIAL_ONLY	full	2/3, -23 to -91	full	2/3, -60 to -228
EQUALIZED_ORDER_SENSITIVITY	full	1/3, -1	full	3/3, -4 to -25
PAIRING_DESTROYED	full	full	full	full

Four findings.

Temporal enrichment alone breaks it. Doubling the temporal resolution while holding the spatial factors fixed leaves 9 of the 21 arm–anchor combinations rank-deficient, with the direct channel deficient at every anchor and losing up to 100 dimensions of 448. Nothing about the geometry changed; only the question being asked of it.

Spatial enrichment is far gentler. At C528_S only 1 of the 21 combinations are deficient — the direct channel at a single anchor, by 6 dimensions of 528. The operator resolves the declared spatial class much better than the declared temporal one — the same asymmetry the mechanism decomposition found from the other side.

At the richest class the full physical operator has a null space. 15 of the 21 combinations are deficient at C1056_ST, including the resolved stack itself (up to 35 dimensions of 1056) and the direct channel (up to 265).

The negative control never loses rank. `PAIRING_DESTROYED` reaches full column rank at every class and every anchor, including the 1056-dimensional one, where the physical resolved operator does not. An operator with no physical content is the only one on the table that stays identifiable under enrichment. This is the same lesson as section 10.1 arriving from the rank side: identifiability is not a measure of physical fidelity.

The delay-only equivalence does not survive enrichment. Indistinguishable from the resolved operator by rank on C224, the two separate once the model is rich enough — and not always in the direction one would guess, because substituting the direct order's well-sampled spatial map onto every order is not a strict impoverishment. `DELAY_ONLY` is a mechanism probe, not a measurement architecture.

11.1 Identifiability and reach move independently

Over the same ladder the resolved *operational*-rank fraction falls by 16.8 percentage points and the smallest determined singular value falls from 1.63e-02 to 1.33e-07 — five orders of magnitude. Recoverable depth moves in 7 of 18 physical anchor–arm rows and never by more than 1 grid step of 4 M.

This is the cleanest statement the campaign supports. **A rank statement is not a depth statement, in either direction, and neither is evidence for the other.** Enriching a source model exposes directions the operator cannot determine; it barely changes how far back in retarded time the operator can see. Papers that report identifiability and papers that report historical reach are reporting different physics, and a result in one does not transfer to the other.

12. Discussion

The picture that survives all of the above is narrower than the one a single-geometry study would support, and we think it is the useful one.

Near-critical null geodesics do create a long, distributed, overlapping retarded-time transfer structure. That structure is not a delay ladder: orders share retarded windows, each ray carries its own delay, and treating the n -th image as a snapshot of one past epoch is an asymptotic convenience that the operator's null structure does not respect.

Higher-order images are drastically suppressed in integrated brightness and retain disproportionate — though not, as we first reported, dramatically disproportionate — Fisher sensitivity to historical source modes. The gap between the throughput and sensitivity exponents is real, roughly a factor of two in the exponent on the interpretable cells, and is largely accounted for by the solid-angle scaling that whitening imposes rather than by anything distinctive about near-critical geodesics.

Recoverable history is controlled by an information-weighted temporal kernel. It is not controlled by the oldest ray, which is an extreme-value statistic that does not converge; not by total photon-ring flux, which the equalized-sensitivity arm shows to be a poor proxy; and not by algebraic rank, which section 11 shows can move by hundreds of dimensions while depth moves by one grid step.

The most transferable methodological point is negative. Three of the quantities one would naturally reach for — total information, algebraic rank, conditioning — each fail as proxies for recoverable history in a specific, demonstrable way: total information is dominated by the bright direct image and barely moves when the informative deep orders are added; rank is a property of the declared class; and conditioning is optimised by a permutation with no physical content at all.

13. What the measurement-model correction changed

The defect and its consequences are recorded in the invalidation ledger as `D-H_flat_sigma_measurement_convention`, and 33 artifacts produced under the retired convention are marked `SUPERSEDED_MEASUREMENT_MODEL_DEFECT`. No table in this manuscript is built from any of them; the builder reads its inputs through the canonical freeze and raises on a superseded path.

What moved, at the canary geometry. The retired-convention values in the middle column are quoted from the defect record and from the pre-correction commit's tables; they are deliberately *not* in the claim ledger, because the ledger may only cite canonical artifacts and those bytes are superseded:

quantity	retired convention	corrected convention
<code>Gamma_sensitivity_matched</code> , 0 to 1	0.576	2.486
<code>Gamma_sensitivity_matched</code> , 1 to 2	0.387	2.120
resolved trace information	—	8.944e+08
direct trace information	—	8.845e+08
resolved-over-direct trace information gain	82%	1.12%
equalized-arm depth advantage	one grid step	tens of M

The first row is the one that mattered scientifically. Under the retired convention the sensitivity exponent was roughly seven to ten times smaller than the throughput exponent, which reads as "information decays far more slowly than flux". Corrected, it is about half — a real effect, but an ordinary one. The 82% trace-information gain becomes 1.12%; the resolved stack collects almost no additional photons, and its value is that its operational rank still rises from 153 to 201 of 224. Structure, not signal.

What did not move under the correction: the ordering of the mechanism decomposition, the negative control's superior conditioning, and full algebraic rank on C224.

The mechanism conclusion was nonetheless weakened, by a separate finding rather than by this correction. Amendment 002 (section 5.2) established that the registered mechanism statistic could not see the delay-only substitution at all. On the diagnostic that can, delay diversity is still the larger contribution but no longer the whole of it. Two distinct problems, corrected independently, both in the direction of a smaller claim.

14. Limitations

1. **Declared classes, not the continuum.** Every rank and nullity is relative to a finite-dimensional class. Section 11 is a demonstration that this qualification is load-bearing, not a formality.
2. **Three image orders.** Orders above $n = 2$ are not computed. No asymptotic exponent in n is fitted, and none should be read off the three points here.
3. **One emissivity and flow prescription.** The throughput exponent is a property of the frozen specific-intensity and Keplerian-flow prescription. It is not the geometric Kerr critical exponent and is not identified with it anywhere.
4. **A source domain that moves with spin.** The primary radial support follows the rays. Section 10.2 measures the confound on three anchors; it does not remove it from the primary surface.
5. **5 matched-support cells are undefined.** At low inclination the $n = 0$ and $n = 1$ retarded windows are disjoint. Those cells are reported and retained in every denominator, but they carry no exponent.
6. **No real data and no learning.** Section 6 inverts synthetic histories only. Nothing is inverted against an observation, and no learned prior, neural estimator or trained model appears anywhere in this campaign.
7. **Single-instrument idealisation.** One monochromatic, spatially resolved intensity observation with white noise of fixed density per unit solid angle. No interferometric sampling, calibration, or polarisation.
8. **The reconstruction result is one geometry and one class.** Section 6 covers $a^* = 0.5$, $i = 50$ degrees and C224. It is not geometry-wide reconstruction and it is not arbitrary movie recovery.
9. **Level, not morphology, at the reference SNR.** 98.4% of a truth's age-window norm is its spatially constant part, and the old-band structure error is 1.169 against 1.141 — above one for both arms. The reconstruction claim is about the age-local emissivity level. The structural result at $SNR_0 = 30000$ is stated separately and must not be merged with it.
10. **Outside the declared class, reconstruction fails.** The off-grid prior-fit regime returns no stable span for any arm, and the tolerance is reachable there for the exact projection, so this is a negative result rather than a non-measurement. The off-grid held-out regime passes but has a much smaller representation floor and is a mild-mismatch diagnostic, not evidence of off-grid robustness.
11. **Uncertainty is withdrawn.** The joint calibration of the probabilistic estimators missed a band frozen in advance, so they are retained as point estimators. No credible interval, posterior movie or coverage statement is available from this line, and none appears in this paper.
12. **The morphology gain is single-feature and heterogeneous.** Section 9 reaches materiality in 5 of 12 family-estimator cells and in no two-feature cell; the absolute assignment cost on stable multi-resolved states stays at 1.208 and 1.302, above one whole

feature wrong. The simplest source in the bank is negative under both estimators.

13. **The morphology baseline is partly floor-limited.** 56.6% of direct-image states under ridge sit at the measure's ceiling, so part of the reduction is states that failed outright under the direct image and did not under the resolved stack. The saturation fractions are reported beside every figure so the reader can discount by inspection.
14. **The morphology bank is smaller than the design it inherited.** It carries 60 truths and 4 paired noise draws where the earlier authorization was 96 and 8. The reduction was fixed and committed before any truth was drawn, so it cannot have shaped the result, but every interval in section 9 is wider than the authorized design would have given.

1. The structural reconstruction test of section 8 is a representation-matched, zero-floor best-case benchmark. Truths are synthesised exactly in the class, which isolates inversion error from representation error and says nothing about arbitrary or realistic accretion-flow histories. An off-class history aligned with well-observed operator directions could be easier than some in-class ones; many others would be much harder.
1. Section 8 reports a preregistered standard that was not met, not a demonstration that the effect is zero. The central estimates are positive on both estimators.
1. The structural effect at the reference SNR lives mainly in a signed constant-flux diagnostic bank, which reaches a negative mass fraction of 0.197 and is not a physical emissivity history. Section 8 reports what excluding it costs.
1. No sealed main was executed for the structural endpoint, so nothing in section 8 is a held-out result. Every sealed commitment is preserved unscored.

Figures. All four are primary; there is no supplementary tier, because a figure not worth placing in the argument is not worth carrying. Each is drawn by `scripts/build_figures.py` from the same canonical tables the claim ledger cites, and each caption is stored with the code that draws it, so a figure and its caption cannot drift apart. Figure 1 is the observability result, figure 2 is the Shiva effect, figure 3 is the two held-out results side by side on separate axes, and figure 4 is the family heterogeneity that qualifies the second of them.

15. Conclusion

A resolved photon ring carries retarded-time information, and this paper is an attempt to say exactly how much, on synthetic data, with the standard for "how much" written down first.

What the orders buy is real and bounded. They see deeper than the direct image at every geometry we registered, by a margin set by inclination and not by spin. They convert an ill-posed old-epoch problem into a well-posed one on compactly supported classes, where the direct image has exact null directions rather than merely poor conditioning. And on two sealed held-out banks they deliver two historical inverse results: `STABLE_BASELINE_INCLUSIVE_EMISSIVITY_LEVEL_RECONSTRUCTION`, an extension of the anchored stable span for age-local emissivity level, and `AGGREGATE_RESOLUTION_AWARE_MORPHOLOGY_ERROR_REDUCTION`, an aggregate improvement in a resolution-aware morphology error, each attributable to resolving the orders rather than to the photons those orders carry.

What they do not buy is a movie. The level result is a level result: most of the recovered norm is the spatially constant part, and old-band structure error stays above one for both arms. The morphology result is an average over a heterogeneous set of source families that reaches materiality in no two-feature cell, extends no stable morphology interval, and is measured against a direct baseline more than half of whose states sit at the measure's ceiling. Age-local structure under our preregistered standard is a bar not cleared at the

reference SNR. We report each of these beside its result rather than after it, because the distance between "a measurable improvement in an error" and "the history was recovered" is where a claim of this kind goes wrong.

The Shiva effect is the reason we think that distance is structural rather than incidental. A source model rich enough to describe a real accretion history is rich enough to destroy the identifiability that made the reach statistic meaningful; a model small enough to be identifiable is a statement about that model. Any future claim to have read history out of a photon ring will have to name which side of that trade it is standing on, and this paper's contribution is to make the trade measurable rather than rhetorical.

Everything here is theory and controlled synthetic computation on one operator family. No telescope detection, no laboratory result, and no recovery from a resolved real photon ring is claimed.

16. Reproducibility and governance

The campaign is gate-driven. Gates carry a mechanical status that is never edited to match a later adjudication, so a defect that was real stays visible:

gates passing:

189

active blocking failures:

0

preserved literal failures:

8

future-phase not run:

11

A preserved literal failure is a FAIL that has been ruled on and kept on the record rather than reinterpreted. There are 8:

preserved failure	adjudicated disposition
EDGE1_a000_i020_raymap_generation	RESOLVED_BY_S0_BACKEND
G10q_retired_flat_sigma_convention	RETIRED_PIXELIZATION_DEPENDENT
G1_v01_reproduction_relative	FAIL_AS_WRITTEN
G7_grid_convergence_a098_i075	RETIRED_NONCONVERGENT_EXTREME_STATISTIC
G7_grid_convergence_raw_max	RETIRED_NONCONVERGENT_EXTREME_STATISTIC
G7b_weighted_operator_discrepancy	WITHDRAWN_INVALID_CONVERGENCE_METRIC
GRID_AUTHORIZATION	SUPERSEDED_GRID_COMPLETE
R0_G15_uncertainty_calibration_band	UNCERTAINTY_WITHDRAWN

Reproduction path: the registry digest is 2ba66f0209fe1cdec97b8cf5862494c22fb94704318f9601a7a1f9eb4b783796 ; the operator grid freeze pins the ray-map hashes, source class, probe, observer sampling, age grid, noise convention, SNR grid, arms, rank conventions, thresholds, censoring rule and permutation seeds before the first geometry was evaluated; the canonical freeze lists the 591 artifacts the manuscript may cite; and the claim ledger maps every number above to the artifact, row filter and column it came from. scripts/verify_manuscript.py re-derives all of them from the frozen bytes and checks that each rendered value appears in this text.

References

Literature entries were checked against the publisher or arXiv record rather than recalled; software entries are pinned to the version or commit this campaign actually ran.

Kerr lensing and the photon ring

1. S. E. Gralla, D. E. Holz and R. M. Wald, *Black hole shadows, photon rings, and lensing rings*, Phys. Rev. D **100**, 024018 (2019), arXiv:1906.00873.
2. S. E. Gralla and A. Lupsasca, *Lensing by Kerr black holes*, Phys. Rev. D **101**, 044031 (2020), arXiv:1910.12873.
3. M. D. Johnson, A. Lupsasca, A. Strominger, G. N. Wong, S. Hadar, D. Kapec, R. Narayan, A. Chael, C. F. Gammie, P. Galison, D. C. M. Palumbo, S. S. Doeleman, L. Blackburn, M. Wielgus, D. W. Pesce, J. R. Farah and J. M. Moran, *Universal interferometric signatures of a black hole's photon ring*, Sci. Adv. **6**, eaaz1310 (2020), arXiv:1907.04329, doi:10.1126/sciadv.aaz1310.
4. Event Horizon Telescope Collaboration, *First M87 Event Horizon Telescope Results. I. The Shadow of the Supermassive Black Hole*, Astrophys. J. Lett. **875**, L1 (2019), doi:10.3847/2041-8213/ab0ec7.

Time domain in the photon ring

1. G. N. Wong, *Black hole glimmer signatures of mass, spin, and inclination*, Astrophys. J. **909**, 217 (2021), arXiv:2009.06641, doi:10.3847/1538-4357/abdd2d.
2. S. Hadar, M. D. Johnson, A. Lupsasca and G. N. Wong, *Photon ring autocorrelations*, Phys. Rev. D **103**, 104038 (2021), arXiv:2010.03683.

References 5 and 6 read the time domain through statistics of the observed intensity — an arrival-time sequence and a two-point function. This paper asks the complementary question: treated as a linear inverse problem, what of the source history is *identifiable and recoverable*, and under what declared class. Neither reference is used as an input to any number here.

Inverse problems

1. P. C. Hansen, *Rank-Deficient and Discrete Ill-Posed Problems: Numerical Aspects of Linear Inversion*, SIAM Monographs on Mathematical Modeling and Computation 4 (1998). The truncated-SVD and Tikhonov estimators used here are the standard ones; nothing in this paper is a new estimator.

Software, pinned to what was run

1. A. Cárdenas-Avendaño, A. Lupsasca and H. Zhu, *Adaptive analytical ray tracing of black hole photon rings*, Phys. Rev. D **107**, 043030 (2023), arXiv:2211.07469. Code: `aart` 2.1.10, the primary ray tracer.
2. A. Chael, `kgeo`, a Kerr null-geodesic tracer implementing the Gralla–Lupsasca formalism, <<https://github.com/achael/kgeo>>, pinned at commit `dc347060c5fb` and used as an independent cross-tracer, not as the primary.
3. NumPy 1.26.4, SciPy 1.11.4, pandas 3.0.5 and PyArrow 25.0.1, under the pinned single-threaded environment described in section 16.

Data and code availability

All code, ray maps, tables, gate records, freezes, amendments and the invalidation ledger are in the campaign repository at tag `paper-I-campaign-final`, commit `8345068676b15ce8f96a76da9d92b159db215f1d`. Ray tracing uses two independent tracers, pinned by version and commit, with an explicitly validated Schwarzschild backend at zero spin.